

Package demo

William Shih, Ricardo Simpao, Nilay Varshney, Luke Yee

3/20/2020

Here is an example of how our package functions run. For our data set, we are using a “SGEMM GPU kernel performance Data Set,” which measures the running times of a matrix-matrix product, given different parameter combinations.

```
library(devtools)
library(tidyverse)
library(STA141CFinal)
library(furrr)

set.seed(141)
dat = read_csv("sgemm_product.csv")
dat = dat[sample(241000, 1000),]
dat2 = dat[1:100,]

#We specify a specific column set
y = dat$`Run1 (ms)`
x = dat[,1:(ncol(dat)-4)]


#linear model objects
fit = linear_reg_bs_C(x, y, s = 10, r = 1000)
```

Linear Regression with blb

```
## Warning: Some expressions had a GC in every iteration; so filtering is disabled.

## # A tibble: 3 x 6
##   expression                                min median `itr/sec`
##   <bch:expr>                                <bch:> <bch:>      <dbl>
## 1 linear_reg_bs(x = x, y = y, s = 10, r = 1000)    22.61s 22.61s    0.0442
## 2 linear_reg_bs_par(x = x, y = y, s = 10, r = 1000) 22.92s 22.92s    0.0436
## 3 linear_reg_bs_C(x, y, s = 10, r = 1000)         2.49s  2.49s    0.401
## # ... with 2 more variables: mem_alloc <bch:byt>, `gc/sec` <dbl>

## Warning in f(...): The default behavior of beeswarm has changed in version
## 0.6.0. In versions <0.6.0, this plot would have been dodged on the y-axis. In
## versions >=0.6.0, grouponX=FALSE must be explicitly set to group on y-axis.
## Please set grouponX=TRUE/FALSE to avoid this warning and ensure proper axis
## choice.
```



95 % Confidence Interval for Variable Coefficients

```
coef_CI(fit, alpha = 0.05)
```

##	Lower_Bounds	Estimates	Upper_Bounds
## Intercept	-250.9593638	-161.822787	-77.756876
## MWG	2.4803281	2.909166	3.349526
## NWG	2.3378785	2.785394	3.245074
## KWG	4.6214054	6.870881	9.220991
## MDIMC	-18.6137757	-15.942214	-13.365062
## NDIMC	-18.0645940	-15.529066	-13.084430
## MDIMA	0.8616486	2.773963	4.808043
## NDIMB	-0.1020447	1.983522	4.069686
## KWI	0.8409772	5.744059	10.709685
## VWM	-1.5499689	9.779875	20.961781
## VWN	-4.8347615	4.594360	14.276281
## STRM	-11.0345567	19.774674	51.951853
## STRN	-57.5970022	-24.403050	9.056510
## SA	0.2486218	28.977774	58.002116
## SB	24.4859861	57.913604	92.355236

```
coef_CI_par(fit,alpha = 0.05)
```

##	Lower_Bounds	Estimates	Upper_Bounds
## Intercept	-250.9593638	-161.822787	-77.756876
## MWG	2.4803281	2.909166	3.349526

```
## NWG      2.3378785    2.785394    3.245074
## KWG      4.6214054    6.870881    9.220991
## MDIMC    -18.6137757 -15.942214 -13.365062
## NDMC     -18.0645940 -15.529066 -13.084430
## MDIMA     0.8616486    2.773963    4.808043
## NDMB     -0.1020447    1.983522    4.069686
## KWI      0.8409772    5.744059   10.709685
## VWM      -1.5499689    9.779875   20.961781
## VWN      -4.8347615    4.594360   14.276281
## STRM     -11.0345567   19.774674   51.951853
## STRN     -57.5970022 -24.403050    9.056510
## SA       0.2486218   28.977774   58.002116
## SB       24.4859861   57.913604   92.355236
```


```
coef_CI_C(fit,alpha = 0.05)
```

```
##           Lower_Bounds  Estimates Upper_Bounds
## Intercept -250.9291528 -161.822787  -76.566608
## MWG        2.4804299    2.909166    3.353727
## NWG        2.3379711    2.785394    3.249154
## KWG        4.6219064    6.870881    9.246523
## MDIMC      -18.6132246 -15.942214 -13.333145
## NDMC       -18.0640532 -15.529066 -13.067617
## MDIMA       0.8620391    2.773963    4.834013
## NDMB       -0.1017314    1.983522    4.088236
## KWI        0.8419045    5.744059   10.754005
## VWM       -1.5468762    9.779875   21.047596
## VWN       -4.8335146    4.594360   14.337845
## STRM      -11.0248991   19.774674   52.324871
## STRN      -57.5908069 -24.403050    9.346352
## SA         0.2563944   28.977774   58.204637
## SB        24.4928834   57.913604   92.451715
```

```
(b1 = bench::mark(
  coef_CI(fit, alpha = 0.05),
  coef_CI_par(fit,alpha = 0.05),
  coef_CI_C(fit, alpha = 0.05),
  check = FALSE)
)
```

```
## # A tibble: 3 x 6
##   expression          min    median `itr/sec` mem_alloc `gc/sec`
##   <bch:expr>        <bch:tm> <bch:tm>    <dbl>   <bch:byt>   <dbl>
## 1 coef_CI(fit, alpha = 0.05)  46.65ms    49ms    19.5     7.49MB    2.17
## 2 coef_CI_par(fit, alpha = 0.05) 83.27ms   106ms    9.75     7.8MB     0
## 3 coef_CI_C(fit, alpha = 0.05)  8.83ms    11ms   76.5     1.15MB     0
```

```
ggplot2::autoplot(b1)
```



Notice that `coef_CI_par` offers better memory allocation than `coef_CI`.

95% Prediction Interval

```
plan(multiprocess, workers = 4)
PI(fit, dat2[1:3, 1:14], alpha = 0.05)
```

```
##      Lower_Bounds Estimates Upper_Bounds
## [1,]    -119.5224   -64.28095    -12.19813
## [2,]     303.9097   373.82120    446.12019
## [3,]    -105.6873   -45.87250     11.50187
```

```
PI_par(fit, dat2[1:3, 1:14], alpha = 0.05)
```

```
##      Lower_Bounds Estimates Upper_Bounds
## [1,]    -119.5224   -64.28095    -12.19813
## [2,]     303.9097   373.82120    446.12019
## [3,]    -105.6873   -45.87250     11.50187
```


```
PI_C(fit, dat2[1:3, 1:14], alpha = 0.05)
```

```
##      Lower_Bounds Estimates Upper_Bounds
## [1,]    -119.5031   -64.28095    -11.66223
## [2,]     303.9193   373.82120    446.62104
## [3,]    -105.6698   -45.87250     12.00876
```

```
(b2 = bench::mark(
  PI(fit, dat2[1:3, 1:14], alpha = 0.05),
  PI_par(fit, dat2[1:3, 1:14], alpha = 0.05),
  PI_C(fit, dat2[1:3, 1:14], alpha = 0.05),
  check = FALSE)
)
```

```
## # A tibble: 3 x 6
##   expression                min   median `itr/sec`
##   <bch:expr>             <bch:tm> <bch:tm>    <dbl>
## 1 PI(fit, dat2[1:3, 1:14], alpha = 0.05)    54.68ms  65.18ms    14.4
## 2 PI_par(fit, dat2[1:3, 1:14], alpha = 0.05) 778.54ms 778.54ms     1.28
## 3 PI_C(fit, dat2[1:3, 1:14], alpha = 0.05)   2.35ms   2.91ms   340.
## # ... with 2 more variables: mem_alloc <bch:byt>, `gc/sec` <dbl>
```

```
ggplot2::autoplot(b2)
```



95 % Confidence Interval for Variance

```
s2_CI(fit, alpha = 0.05)
```

```
## Lower_Bound   Estimate Upper_Bound
##    572738.9    718571.7   858566.8
```

```
s2_CI_par(fit, alpha = 0.05)
```

```
## Lower_Bound Estimate Upper_Bound
## 572738.9 718571.7 858566.8
```


```
s2_CI_C(fit, alpha = 0.05)
```

```
## Lower_Bound Estimate Upper_Bound
## 572781.4 718571.7 859655.6
```

```
(b3 = bench::mark(
  s2_CI(fit, alpha = 0.05),
  s2_CI_par(fit, alpha = 0.05),
  s2_CI_C(fit, alpha = 0.05),
  check = FALSE)
)
```

```
## # A tibble: 3 x 6
## expression min median `itr/sec` mem_alloc `gc/sec`
## <bch:expr> <bch:tm> <bch:tm> <dbl> <bch:byt> <dbl>
## 1 s2_CI(fit, alpha = 0.05) 1.67ms 1.92ms 461. 118.12KB 0
## 2 s2_CI_par(fit, alpha = 0.05) 97.81ms 102.18ms 9.56 2.53MB 2.39
## 3 s2_CI_C(fit, alpha = 0.05) 97.4us 98.6us 8910. 2.49KB 0
```

```
ggplot2::autoplot(b3)
```



Notice that `s2_CI_par` offers better memory allocation than `s2_CI`.